

MLW / KUHeS Statistics and R short course

Session 1 - Practical (solutions)

Marc Henrion, Evaristar Kudowa, Vester Gunsaru

2024-10-08

Session 1 - Practical (Solutions)

Go to the course website on GitHub:

https://github.com/mlw-stats/R_And_Statistics_Training_Autumn2024/Session1

From here, download the following files:

btTBreg.csv

btTBregHospitals.csv

btTBreg_info.txt

1. Load the `btTBreg.csv` data table into R.

```
btDat<-read.csv("dataAndSupportDocs/btTBreg.csv")
```

```
head(btDat) # have a look at the data
```

	id	age	sex	hiv	bmi	ses	cd41	cd42	cd41.sk	cd42.sk	hosp
1	1	44	2	0	26.32	4	346	519	313.11656	572.8906	1
2	2	32	2	0	20.79	5	237	337	43.12752	406.1971	5
3	3	32	1	0	19.21	1	198	328	338.32172	408.2427	2
4	4	20	1	0	21.34	4	246	525	77.08697	312.7572	3
5	5	30	1	0	23.98	4	270	444	169.02539	335.3739	3
6	6	32	1	0	17.97	4	283	372	255.45773	323.4773	4

```
dim(btDat) # check dimesnions of data table
```

[1] 3000 11

2. The variables `cd41`, `cd42` and `cd41.sk`, `cd42.sk` measure the same variables (`cd4` and `cd4.sk` respectively) in the same individuals at two different time point. This means the data are in wide format. Reformat to long format.

The key difficulty here is that you have 2 variables (at 2 times points). In the example from lectures we only had 1 variable (at 3 different conditions). One approach is to do each variable separately, then combine the resulting data frames:

```
btDatLong.cd4<-btDat %>%
  pivot_longer(names_to="time", values_to="cd4", cols=c(cd41, cd42)) %>%
  select(id,age,sex,hiv,bmi,ses,hosp,time,cd4)

btDatLong.cd4sk<-btDat%>%
  pivot_longer(names_to="time", values_to="cd4.sk", cols=c(cd41.sk, cd42.sk)) %>%
  select(id,age,sex,hiv,bmi,ses,hosp,time,cd4.sk)

btDatLong<-data.frame(btDatLong.cd4,cd4.sk=btDatLong.cd4sk$cd4.sk)

rm(btDatLong.cd4,btDatLong.cd4sk)

btDatLong$time<-factor(
  case_when(
    btDatLong$time=="cd41"~"entry",
    btDatLong$time=="cd42"~"exit",
    TRUE~NA_character_)
  ) # rename the levels of the time variable

head(btDatLong) # have a look at the data
```

	id	age	sex	hiv	bmi	ses	hosp	time	cd4	cd4.sk
1	1	44	2	0	26.32	4	1	entry	346	313.11656
2	1	44	2	0	26.32	4	1	exit	519	572.89062
3	2	32	2	0	20.79	5	5	entry	237	43.12752
4	2	32	2	0	20.79	5	5	exit	337	406.19707
5	3	32	1	0	19.21	1	2	entry	198	338.32172
6	3	32	1	0	19.21	1	2	exit	328	408.24267

```
dim(btDatLong) # check dimensions
```

```
[1] 6000    10
```

This can be done a bit more directly, by using regular expression (character expressions that match flexibly to names) and a combination of `pivot_longer()` and `pivot_wider()`:

```
btDatLong<-btDat %>%
  pivot_longer(cols=c(cd41, cd42, cd41.sk, cd42.sk),
               names_pattern = "cd4(1|2)(.*)",
               names_to = c("time","cd4"),
               values_to="measurement") %>%
  mutate(cd4=paste(sep="", "cd4", cd4)) %>%
  pivot_wider(names_from=cd4, values_from=measurement)

head(btDatLong) # have a look at the data
```

```
# A tibble: 6 x 10
```

	id	age	sex	hiv	bmi	ses	hosp	time	cd4	cd4.sk
	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>	<chr>	<dbl>	<dbl>
1	1	44	2	0	26.3	4	1	1	346	313.
2	1	44	2	0	26.3	4	1	2	519	573.
3	2	32	2	0	20.8	5	5	1	237	43.1
4	2	32	2	0	20.8	5	5	2	337	406.
5	3	32	1	0	19.2	1	2	1	198	338.
6	3	32	1	0	19.2	1	2	2	328	408.

```
dim(btDatLong) # check dimensions
```

```
[1] 6000    10
```

The code above requires a bit of unpacking:

- The expression in brackets in the “names_pattern” argument are regular expression matching sequences of character: “(1|2)” matches 1 or 2 and “(.)” matches anything.
- The “mutate()” line is needed as the values store in the “cd4” column are “ ” and “.sk” – as the `pivot_wider()` statement on the next line will use those as column names, we cannot have an empty column name – “ ” would trigger an error message. So we just add the characters “cd4” in front of the stores values – i.e. we then have “cd4” and “cd4.sk” rather than “ ” and “.sk”.

An alternative function that can be used is `reshape()`. To get more information on this function, type `?reshape` at the console.

```
btDatLong<-reshape(btDat,
                    direction="long",
                    varying=list(c("cd41","cd42"),c("cd41.sk","cd42.sk")),
                    ids="id",
                    v.names=c("cd4","cd4.sk"))
```

```
head(btDatLong) # have a look at the data
```

	id	age	sex	hiv	bmi	ses	hosp	time	cd4	cd4.sk
1.1	1	44	2	0	26.32	4	1	1	346	313.11656
2.1	2	32	2	0	20.79	5	5	1	237	43.12752
3.1	3	32	1	0	19.21	1	2	1	198	338.32172
4.1	4	20	1	0	21.34	4	3	1	246	77.08697
5.1	5	30	1	0	23.98	4	3	1	270	169.02539
6.1	6	32	1	0	17.97	4	4	1	283	255.45773

```
dim(btDatLong) # check dimensions
```

```
[1] 6000    10
```

3. Save the reformatted data into a file called `btTBregLong.tab` in such a way that
 - i. Columns are tab-separated.
 - ii. Column names are saved.
 - iii. No row number is saved in the resulting file.

```
dir.create("Session1_output",showWarnings=F)
write.table(btDatLong,sep="\t",col.names=T,row.names=F,file="Session1_output/btTBregLong.t
```

4. Copy the code below to generate some wide-format data. We will assume this dataset contains observations of 2 biomarkers, `ferritin` and `rbp4` for 10 study participants at 2 different timepoints, `day1` and `day90`.

```
set.seed(123)
```

```
df<-data.frame(
  id=paste(sep="", "P", 1:10),
  ferritin_day1=rexp(10,rate=1/195),
  rbp4_day1=rexp(10,rate=1/2.5)
```

```

) %>%
  mutate(
    ferritin_day90=rnorm(10,mean=ferritin_day1+5,sd=4),
    rbp4_day90=rbp4_day1+rexp(10,rate=1/0.25)
  )

```

This is what this data table looks like:

	id	ferritin_day1	rbp4_day1	ferritin_day90	rbp4_day90
1	P1	164.474166	2.5120751	169.251422	3.021712
2	P2	112.439003	1.2005368	114.301474	1.552174
3	P3	259.165699	0.7025341	261.231686	0.829188
4	P4	6.157585	0.9427946	10.294123	1.007684
5	P5	10.961140	0.4707101	14.621489	1.119933
6	P6	61.717737	2.1244653	62.374941	2.431722
7	P7	61.274322	3.9080088	65.932629	4.105679
8	P8	28.327027	1.1969010	37.609469	1.354221
9	P9	531.616111	1.4773371	536.034536	1.790997
10	P10	5.684922	10.1025293	6.022743	10.249700

Reformat this to long format, i.e. so that you have 4 columns: id, time, ferritin and rbp4.

```

dfLong<-df %>%
  pivot_longer(cols=c(ferritin_day1,rbp4_day1,ferritin_day90,rbp4_day90),
    names_pattern="(.)_(.)",
    names_to=c("biomarker","timepoint"),
    values_to="value") %>%
  pivot_wider(names_from=biomarker,values_from=value)

print(dfLong)

```

A tibble: 20 x 4

	id	timepoint	ferritin	rbp4
	<chr>	<chr>	<dbl>	<dbl>
1	P1	day1	164.	2.51
2	P1	day90	169.	3.02
3	P2	day1	112.	1.20
4	P2	day90	114.	1.55
5	P3	day1	259.	0.703
6	P3	day90	261.	0.829

```

7 P4    day1      6.16  0.943
8 P4    day90     10.3   1.01
9 P5    day1     11.0   0.471
10 P5   day90     14.6   1.12
11 P6   day1     61.7   2.12
12 P6   day90     62.4   2.43
13 P7   day1     61.3   3.91
14 P7   day90     65.9   4.11
15 P8   day1     28.3   1.20
16 P8   day90     37.6   1.35
17 P9   day1    532.    1.48
18 P9   day90    536.    1.79
19 P10  day1      5.68 10.1
20 P10  day90      6.02 10.2

```

5. Load the `btTBregHospitals.csv` data table. Join the data frames storing `btTBreg.csv` and `btTBregHospitals.csv`.

```

btDatHosp<-read.csv("dataAndSupportDocs/btTBregHospitals.csv")

head(btDatHosp) # have a look at the data

```

	HID	ShortName	FullName	beds	city
1	1	QECH Queen Elizabeth Central Hospital	1000	Blantyre	
2	2	KCH Kamuzu Central Hospital	1000	Lilongwe	
3	3	ZCH Zomba Central Hospital	400	Zomba	
4	4	MCH Mzuzu Central Hospital	350	Mzuzu	
5	5	Mlambe Mission Hospital	254	Lunzu	

```

dim(btDatHosp) # check dimensions of the data table

```

```

[1] 5 5

```

```

btDatJoined<-btDat %>%
  inner_join(btDatHosp,by=c("hosp"="HID"))

head(btDatJoined) # have a look

```

	id	age	sex	hiv	bmi	ses	cd41	cd42	cd41.sk	cd42.sk	hosp	ShortName
1	1	44	2	0	26.32	4	346	519	313.11656	572.8906	1	QECH
2	2	32	2	0	20.79	5	237	337	43.12752	406.1971	5	Mlambe
3	3	32	1	0	19.21	1	198	328	338.32172	408.2427	2	KCH
4	4	20	1	0	21.34	4	246	525	77.08697	312.7572	3	ZCH
5	5	30	1	0	23.98	4	270	444	169.02539	335.3739	3	ZCH
6	6	32	1	0	17.97	4	283	372	255.45773	323.4773	4	MCH

	FullName	beds	city
1	Queen Elizabeth Central Hospital	1000	Blantyre
2	Mlambe Mission Hospital	254	Lunzu
3	Kamuzu Central Hospital	1000	Lilongwe
4	Zomba Central Hospital	400	Zomba
5	Zomba Central Hospital	400	Zomba
6	Mzuzu Central Hospital	350	Mzuzu

```
dim(btDatJoined) # check dimensions
```

```
[1] 3000 15
```

6. Compute the average patient age and the proportion of male patients for each hospital.

Useful functions for this are `aggregate()` and `group_by()`. You can however also do it manually.

- Manually:

```
# initialise new variables
btDatHosp$avgAge<-NA
btDatHosp$propMale<-NA

# iterate over hospitals
for(i in 1:nrow(btDatHosp)){
  btDatHosp$avgAge[i]<-mean(btDatJoined$age[btDatJoined$ShortName==btDatHosp$ShortName[i]])
  btDatHosp$propMale[i]<-sum(btDatJoined$sex==1 &
                           btDatJoined$ShortName==btDatHosp$ShortName[i]) /
                           sum(btDatJoined$ShortName==btDatHosp$ShortName[i])
}

print(btDatHosp)
```

	HID	ShortName	FullName	beds	city	avgAge
1	1	QECH	Queen Elizabeth Central Hospital	1000	Blantyre	33.14020
2	2	KCH	Kamuzu Central Hospital	1000	Lilongwe	32.80067
3	3	ZCH	Zomba Central Hospital	400	Zomba	32.99310
4	4	MCH	Mzuzu Central Hospital	350	Mzuzu	32.87382
5	5	Mlambe	Mlambe Mission Hospital	254	Lunzu	32.89950

	propMale
1	0.4763514
2	0.4757119
3	0.4948276
4	0.4731861
5	0.5242881

- Using `aggregate()`

```
btDat$hosp<-factor(btDat$hosp)
btDatHosp$avgAge<-aggregate(btDatJoined$age,FUN=mean,by=list(btDat$hosp))$x
btDatHosp$propMale<-aggregate(ifelse(btDatJoined$sex==1,1,0),FUN=mean,by=list(btDat$hosp))

print(btDatHosp)
```

	HID	ShortName	FullName	beds	city	avgAge
1	1	QECH	Queen Elizabeth Central Hospital	1000	Blantyre	33.14020
2	2	KCH	Kamuzu Central Hospital	1000	Lilongwe	32.80067
3	3	ZCH	Zomba Central Hospital	400	Zomba	32.99310
4	4	MCH	Mzuzu Central Hospital	350	Mzuzu	32.87382
5	5	Mlambe	Mlambe Mission Hospital	254	Lunzu	32.89950

	propMale
1	0.4763514
2	0.4757119
3	0.4948276
4	0.4731861
5	0.5242881

- Using `group_by()`

```
tmp<-btDat %>%
  group_by(hosp) %>%
  summarise(avgAge=mean(age,na.rm=T))
btDatHosp$avgAge<-tmp$avgAge
```



```

tmp<-btDat %>%
  group_by(hosp) %>%
  summarise(propMale=mean(ifelse(sex==1,1,0),na.rm=T))
btDatHosp$propMale<-tmp$propMale

print(btDatHosp)

```

	HID	ShortName	FullName	beds	city	avgAge
1	1	QECH	Queen Elizabeth Central Hospital	1000	Blantyre	33.14020
2	2	KCH	Kamuzu Central Hospital	1000	Lilongwe	32.80067
3	3	ZCH	Zomba Central Hospital	400	Zomba	32.99310
4	4	MCH	Mzuzu Central Hospital	350	Mzuzu	32.87382
5	5	Mlambe	Mlambe Mission Hospital	254	Lunzu	32.89950

	propMale
1	0.4763514
2	0.4757119
3	0.4948276
4	0.4731861
5	0.5242881

7. Write an R function that computes the following summary statistics, then, using your custom function, compute these for the bmi, cd41, cd42 columns:

- i. mean
- ii. median
- iii. interquartile range
- iv. minimum
- v. maximum
- vi. number of missing values

```

summaryFun<-function(x){
  return(c(
    mean(x,na.rm=T),
    median(x),
    paste(sep="","(",paste(collapse=" ",quantile(x,probs=c(0.25,0.75))),"),"),
    min(x,na.rm=T),
    max(x,na.rm=T),
    sum(is.na(x))
  ))
}

res<-apply(btDat[,c("bmi","cd41","cd42")],MARGIN=2,FUN=summaryFun)

```

```
rownames(res)<-c("mean","median","IQR","min","max","num_MV")
print(res)
```

	bmi	cd41	cd42
mean	"23.05743333333333"	"248.7943333333333"	"448.003"
median	"23.05"	"249"	"447"
IQR	"(21.34,24.74)"	"(216,281)"	"(381,515)"
min	"12.64"	"57"	"81"
max	"31.14"	"447"	"843"
num_MV	"0"	"0"	"0"

8. Do the same now, but only for female patients. Repeat for only male patients.

```
resF<-apply(btDat[btDat$sex==2,c("bmi","cd41","cd42")],MARGIN=2,FUN=summaryFun)
rownames(resF)<-c("mean","median","IQR","min","max","num_MV")
print(resF)
```

	bmi	cd41	cd42
mean	"23.1218644067797"	"248.473924380704"	"446.675358539765"
median	"23.14"	"250"	"447.5"
IQR	"(21.365,24.82)"	"(215,281)"	"(379,512)"
min	"12.64"	"57"	"138"
max	"31.14"	"447"	"820"
num_MV	"0"	"0"	"0"

```
resM<-apply(btDat[btDat$sex==1,c("bmi","cd41","cd42")],MARGIN=2,FUN=summaryFun)
rownames(resM)<-c("mean","median","IQR","min","max","num_MV")
print(resM)
```

	bmi	cd41	cd42
mean	"22.9900136425648"	"249.129604365621"	"449.392223738063"
median	"22.98"	"248"	"447"
IQR	"(21.3,24.66)"	"(216,282)"	"(383,519.75)"
min	"14.44"	"71"	"81"
max	"30.9"	"414"	"843"
num_MV	"0"	"0"	"0"